



DATA MANIPULATION IN R WITH DATA.TABLE

Welcome to the course!

Matt Dowle and Arun Srinivasan
Instructors, DataCamp



What is a data.table?

- Enhanced `data.frame`
 - inherits from and extends `data.frame`
- Columnar data structure
- Every column must be of same length but can be of different type



Why use data.table?

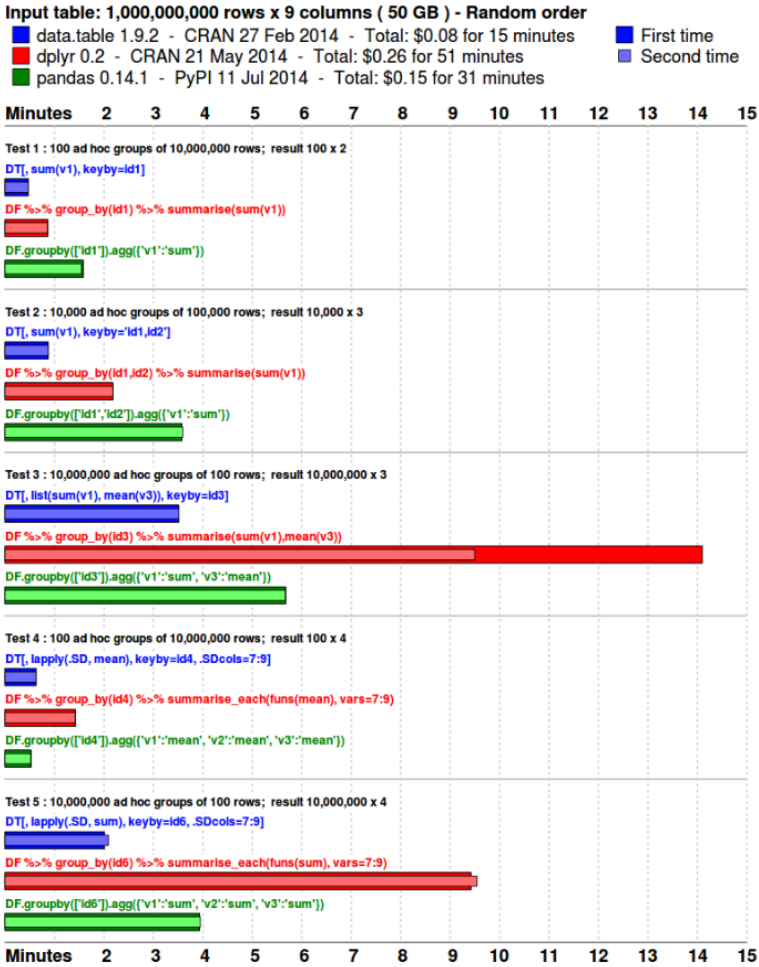
- Concise and consistent syntax
 - Think in terms of `rows`, `columns` and `groups`
 - provides a *placeholder* for each

```
# General form of data.table syntax
DT[i, j, by]
  |  |  |
  |  |  --> grouped by what?
  |  -----> what to do?
  -----> on which rows?
```



Why use data.table?

Fast and memory efficient





Why use data.table?

- Feature-rich
 - Parallelisation
 - Fast updates *by reference*
 - Powerful joins ([Joining Data in R with data.table](#))



Creating a data.table (I)

Three ways of creating data tables:

- `data.table()`
- `as.data.table()`
- `fread()`



Creating a data.table (II)

```
library(data.table)

x_df <- data.frame(id = 1:2,
                  name = c("a", "b"))

x_df
  id name
1  1   a
2  2   b

x_dt <- data.table(id = 1:2,
                  name = c("a", "b"))

x_dt
  id name
1:  1   a
2:  2   b
```



Creating a data.table (III)

```
y <- list(id = 1:2, name = c("a", "b"))
```

```
y
```

```
$id
```

```
[1] 1 2
```

```
$name
```

```
[1] "a" "b"
```

```
x <- as.data.table(y)
```

```
x
```

```
   id name
```

```
1:  1    a
```

```
2:  2    b
```




data.tables and data.frames (I)

Since a data.table *is* a data.frame ...

```
x <- data.table(id = 1:2,  
               name = c("a", "b"))
```

```
x  
   id name  
1:  1    a  
2:  2    b
```

```
class(x)  
[1] "data.table" "data.frame"
```



data.tables and data.frames (II)

Functions used to query data.frames also work on data.tables

```
nrow(x)
[1] 2

ncol(x)
[1] 2

dim(x)
[1] 2 2
```

data.tables and data.frames (III)

data table never automatically converts character columns to factors

```
x_df <- data.frame(id = 1:2,  
                  name = c("a", "b"))  
class(x_df$name)  
[1] "factor"  
  
x_dt <- data.table(id = 1:2,  
                  name = c("a", "b"))  
class(x_dt$name)  
[1] "character"
```



data.tables and data.frames (IV)

Never sets, needs or uses *row names*

```
rownames(x_dt) <- c("R1", "R2")
```

```
x_dt
```

```
  id name
```

```
1:  1    a
```

```
2:  2    b
```



DATA MANIPULATION IN R WITH DATA.TABLE

Let's practice!



DATA MANIPULATION IN R WITH DATA.TABLE

Filtering rows in a data.table

Matt Dowle and Arun Srinivasan
Instructors, DataCamp



General form of data.table syntax

First argument `i` is used to *subset* or *filter* rows

```
# General form of data.table syntax
DT[i, j, by]
  |  |  |
  |  |  --> grouped by what?
  |  -----> what to do?
  -----> on which rows?
```



Row numbers

```
# Subset 3rd and 4th rows from batrips  
batrips[3:4]
```

```
# Same as  
batrips[3:4, ]
```

```
# Subset everything except first five rows  
batrips[-(1:5)]
```

```
# Same as  
batrips[!(1:5)]
```


Special symbol .N

- `.N` is an integer value that contains the number of rows in the `data.table`
- Particularly useful alternative to `nrow(x)` in `i`

```
nrow(batrips)
[1] 326339

batrips[326339]
  trip_id duration ...
1:  588914      364 ...

# Returns the last row
batrips[.N]
  trip_id duration ...
1:  588914      364 ...

# Return all but the last 10 rows
ans <- batrips[1:(.N-10)]
nrow(ans)
[1] 326329
```



Logical expressions (I)

```
# Subset rows where subscription_type is "Subscriber"
batrips[subscription_type == "Subscriber"]

# If batrips was only a data frame
batrips[batrips$subscription_type == "Subscriber", ]
```



Logical expressions (II)

```
# Subset rows where start_terminal = 58 and end_terminal is not 65
batrips[start_terminal == 58 & end_terminal != 65]

# If batrips was only a data frame
batrips[batrips$start_terminal == 58 & batrips$end_terminal != 65]
```



Logical expressions (III)

Optimized using secondary indices for speed automatically.

```
set.seed(1)
dt <- data.table(x = sample(10000, 10e6, TRUE),
                 y = sample(letters, 1e6, TRUE))

indices(dt)
NULL

# 0.207s on first run (time to create index + subset)
system.time(dt[x == 900])
user  system elapsed
0.207   0.015   0.226

indices(dt)
[1] "x"

# 0.002s on subsequent runs (instant subset using index)
system.time(dt[x == 900])
user  system elapsed
0.002   0.000   0.002
```



DATA MANIPULATION IN R WITH DATA.TABLE

Let's practice!



DATA MANIPULATION IN R WITH DATA.TABLE

Helpers for filtering

Matt Dowle and Arun Srinivasan
Instructors, DataCamp



%like%

- `%like%` allows you to search for a *pattern* in a *character* or a *factor* vector
 - Usage: `col %like% pattern`

```
# Subset all rows where start_station starts with San Francisco
batrips[start_station %like% "^San Francisco"]

# Instead of
batrips[grepl("^San Francisco", start_station)]
```



%between%

- `%between%` allows you to search for values in the closed interval `[val1, val2]`
 - **Usage:** `numeric_col %between% c(val1, val2)`

```
# Subset all rows where duration is between 2000 and 3000
batrips[duration %between% c(2000, 3000)]
```

```
# Instead of
batrips[duration >= 2000 & duration <= 3000]
```


%chin%

- `%chin%` is similar to `%in%`, but it is *much* faster and only for character vectors
 - **Usage:** `character_col %chin% c("val1", "val2", "val3")`

```
# Subset all rows where start_station is  
# "Japantown", "Mezes Park" or "MLK Library"  
batrips[start_station %chin% c("Japantown", "Mezes Park", "MLK Library")]  
  
# much faster than  
batrips[start_station %in% c("Japantown", "Mezes Park", "MLK Library")]
```



DATA MANIPULATION IN R WITH DATA.TABLE

Let's practice!